

## Assignment 1: Write a program to implementing and evaluating a Linear Regression model

```
# Import necessary libraries
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error, r2_score

# Load the dataset from a CSV file
data = pd.read_csv('Data science II/advertising.csv')
# Check the first few rows of the dataset to understand its structure
print(data)

# Define the independent variable (feature) and dependent variable (target)
X = data[['TV']] # Independent variable (1D array, needs to be 2D for sklearn)
y = data['Sales'] # Dependent variable (target)

# Split the data into training and testing sets (80% train, 20% test)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Initialize and train the Linear Regression model
model = LinearRegression()
model.fit(X_train, y_train)

# Make predictions on the test set
y_pred = model.predict(X_test)

# Evaluate the model's performance
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)

# Output the evaluation metrics
print(f'Mean Squared Error (MSE): {mse}')
print(f'R-squared (R2) Score: {r2}')

# Visualize the results
plt.scatter(X_test, y_test, color='blue', label='True Values')
```


```

plt.plot(X_test, y_pred, color='red', label='Regression Line')
plt.xlabel('TV')
plt.ylabel('Sales')
plt.title('Linear Regression Model')
plt.legend()
plt.show()

```

## Output:

Mean Squared Error (MSE): 6.101072906773964  
R-squared (R2) Score: 0.802561303423698



**Assignment 2: Write a program to implementing and evaluating a Logistic Regression model.**

```

import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression

```

```

from sklearn.metrics import accuracy_score, classification_report, confusion_matrix

# Load dataset from CSV file
df = pd.read_csv('log.csv')

# Assuming the last column is the target variable
X = df.iloc[:, :-1]
y = df.iloc[:, -1]

# Split dataset into train and test sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Standardize features
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)

# Train Logistic Regression model
model = LogisticRegression()
model.fit(X_train, y_train)

# Make predictions
y_pred = model.predict(X_test)

# Evaluate the model
accuracy = accuracy_score(y_test, y_pred)
conf_matrix = confusion_matrix(y_test, y_pred)
report = classification_report(y_test, y_pred)

# Print evaluation results
print(f'Accuracy: {accuracy:.4f}')
print('Confusion Matrix:')
print(conf_matrix)
print('Classification Report:')
print(report)

```

## Output:

```
Accuracy: 0.4650
Confusion Matrix:
[[46 50]
 [57 47]]
Classification Report:
```

	precision	recall	f1-score	support
0	0.45	0.48	0.46	96
1	0.48	0.45	0.47	104
accuracy			0.47	200
macro avg	0.47	0.47	0.46	200
weighted avg	0.47	0.47	0.47	200

## Assignment 3: Write a program to implementing and evaluating a Decision Tree classifier.

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.tree import DecisionTreeClassifier, plot_tree
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix

# Load dataset from CSV file
df = pd.read_csv('log.csv')

# Assuming the last column is the target variable
X = df.iloc[:, :-1]
y = df.iloc[:, -1]

# Split dataset into train and test sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Standardize features (optional for Decision Tree, but can help with performance)
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
```

```

# Train Decision Tree model
model = DecisionTreeClassifier()
model.fit(X_train, y_train)

# Make predictions
y_pred = model.predict(X_test)

# Evaluate the model
accuracy = accuracy_score(y_test, y_pred)
conf_matrix = confusion_matrix(y_test, y_pred)
report = classification_report(y_test, y_pred)

# Print evaluation results
print(f'Accuracy: {accuracy:.4f}')
print('Confusion Matrix:')
print(conf_matrix)
print('Classification Report:')
print(report)

# Visualize the Decision Tree
plt.figure(figsize=(15, 10))
plot_tree(model, filled=True, feature_names=df.columns[:-1], class_names=str(np.unique(y)))
plt.show()

```

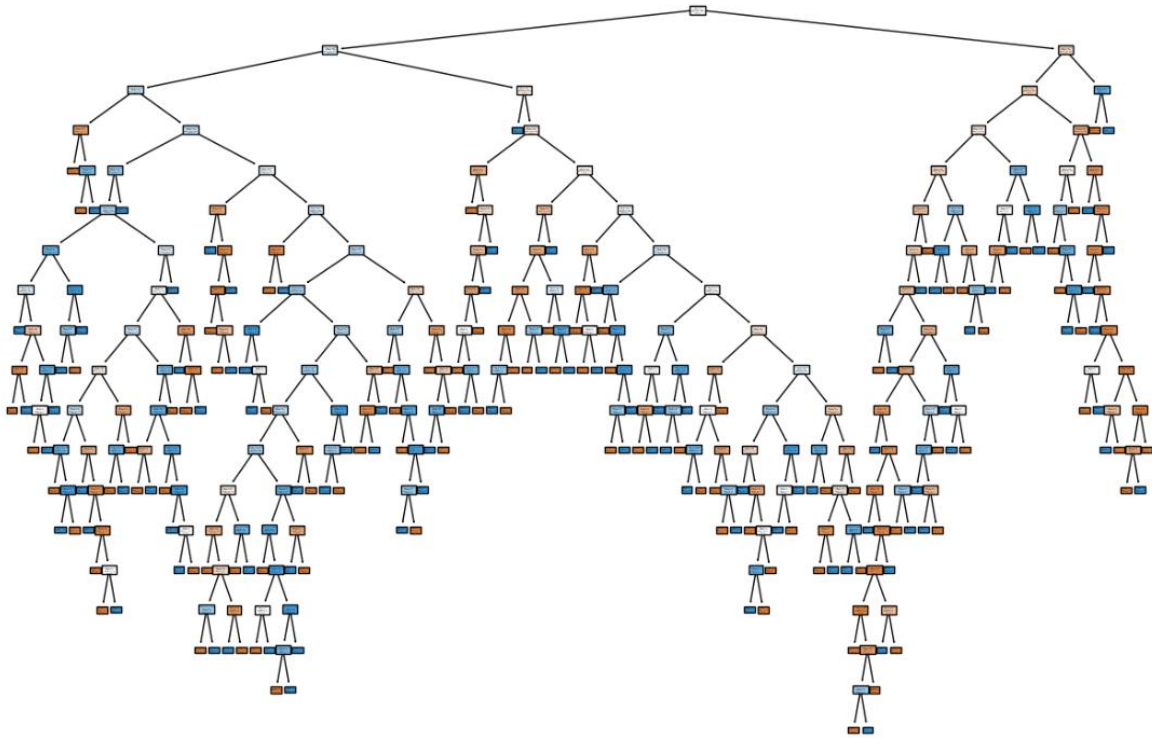
### Output:

```

Accuracy: 0.4650
Confusion Matrix:
[[45 51]
 [56 48]]
Classification Report:

```

	precision	recall	f1-score	support
0	0.45	0.47	0.46	96
1	0.48	0.46	0.47	104
accuracy			0.47	200
macro avg	0.47	0.47	0.46	200
weighted avg	0.47	0.47	0.47	200



#### Assignment 4: Write a program to implementing Clustering using the K-means algorithm

```
import numpy as np
import matplotlib.pyplot as plt
from sklearn.cluster import KMeans
from sklearn.datasets import make_blobs

# Step 1: Generate synthetic data (for demonstration)
# Generating 300 data points with 4 centers (clusters)
X, y = make_blobs(n_samples=300, centers=4, random_state=42)

# Step 2: Apply the K-means algorithm
# Set the number of clusters to 4 (since we generated data with 4 centers)
kmeans = KMeans(n_clusters=4, random_state=42)
kmeans.fit(X)

# Step 3: Get the centroids and labels for the clusters
centroids = kmeans.cluster_centers_
labels = kmeans.labels_
```

```
# Step 4: Visualize the clusters
```

```
plt.figure(figsize=(8, 6))
```

```
# Scatter plot of data points with colors corresponding to cluster labels
```

```
plt.scatter(X[:, 0], X[:, 1], c=labels, cmap='viridis', marker='o', edgecolor='k')
```

```
# Mark the centroids with a red 'X'
```

```
plt.scatter(centroids[:, 0], centroids[:, 1], marker='X', s=200, c='red', label='Centroids')
```

```
# Add titles and labels
```

```
plt.title('K-means Clustering')
```

```
plt.xlabel('Feature 1')
```

```
plt.ylabel('Feature 2')
```


```
# Show the legend
```

```
plt.legend()
```

```
# Display the plot
```

```
plt.show()
```

**Output:**



### Assignment 5: Write a program to implementing Dimensionality reduction using PCA.

```
# Import necessary libraries
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.decomposition import PCA
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split

# Load CSV file (replace 'your_dataset.csv' with your actual dataset file path)
df = pd.read_csv('log.csv')

# Check the first few rows of the dataset
print(df.head())

# Step 1: Separate features (X) and target (y) if applicable
# Assuming the last column is the target variable
X = df.iloc[:, :-1].values # All rows, all columns except the last one
y = df.iloc[:, -1].values # Last column is the target

# Step 2: Standardize the dataset (important for PCA)
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

# Step 3: Apply PCA to reduce to 2 dimensions for visualization
pca = PCA(n_components=2) # Reduce to 2 components for visualization
X_pca = pca.fit_transform(X_scaled)

# Step 4: Explained variance ratio (how much variance is captured by each component)
print("Explained variance ratio:", pca.explained_variance_ratio_)

# Step 5: Plot the 2D PCA result
plt.figure(figsize=(8,6))
plt.scatter(X_pca[:, 0], X_pca[:, 1], c=y, cmap='viridis', edgecolor='k', s=100)
plt.title("PCA of Dataset")
plt.xlabel("Principal Component 1")
plt.ylabel("Principal Component 2")
plt.colorbar(label="Target Class")
plt.show()
```




Output:

	feature_0	feature_1	feature_2	feature_3	feature_4	feature_5	\
0	0.496714	-0.138264	0.647689	1.523030	-0.234153	-0.234137	
1	-0.463418	-0.465730	0.241962	-1.913280	-1.724918	-0.562288	
2	1.465649	-0.225776	0.067528	-1.424748	-0.544383	0.110923	
3	-0.601707	1.852278	-0.013497	-1.057711	0.822545	-1.220844	
4	0.738467	0.171368	-0.115648	-0.301104	-1.478522	-0.719844	

	feature_6	feature_7	feature_8	feature_9	target
0	1.579213	0.767435	-0.469474	0.542560	1
1	-1.012831	0.314247	-0.908024	-1.412304	1
2	-1.150994	0.375698	-0.600639	-0.291694	0
3	0.208864	-1.959670	-1.328186	0.196861	0
4	-0.460639	1.057122	0.343618	-1.763040	0

Explained variance ratio: [0.11348263 0.10726962]



## Assignment 6: Write a program to implementing Bagging using Random Forest.

```
# Import necessary libraries
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score
from sklearn.preprocessing import LabelEncoder

# Load CSV file (replace 'your_dataset.csv' with your actual dataset file path)
df = pd.read_csv('log.csv')

# Check the first few rows of the dataset
print(df.head())

# Step 1: Handle missing values (if any)
# Example: Drop rows with missing values (you can also fill with the mean or median)
df = df.dropna()

# Step 2: Separate features (X) and target (y)
# Assuming the last column is the target variable
X = df.iloc[:, :-1].values # All rows, all columns except the last one (features)
y = df.iloc[:, -1].values # Last column is the target

# Step 3: Encode the target variable if it's categorical
label_encoder = LabelEncoder()
y = label_encoder.fit_transform(y)

# Step 4: Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)

# Step 5: Initialize and train the Random Forest Classifier
rf_classifier = RandomForestClassifier(n_estimators=100, random_state=42)
rf_classifier.fit(X_train, y_train)

# Step 6: Make predictions on the test set
y_pred = rf_classifier.predict(X_test)

# Step 7: Evaluate the model's performance
accuracy = accuracy_score(y_test, y_pred)
```

```
print(f'Accuracy of Random Forest on test data: {accuracy * 100:.2f}%')
```

**Output :**

	feature_0	feature_1	feature_2	feature_3	feature_4	feature_5	\
0	0.496714	-0.138264	0.647689	1.523030	-0.234153	-0.234137	
1	-0.463418	-0.465730	0.241962	-1.913280	-1.724918	-0.562288	
2	1.465649	-0.225776	0.067528	-1.424748	-0.544383	0.110923	
3	-0.601707	1.852278	-0.013497	-1.057711	0.822545	-1.220844	
4	0.738467	0.171368	-0.115648	-0.301104	-1.478522	-0.719844	

	feature_6	feature_7	feature_8	feature_9	target
0	1.579213	0.767435	-0.469474	0.542560	1
1	-1.012831	0.314247	-0.908024	-1.412304	1
2	-1.150994	0.375698	-0.600639	-0.291694	0
3	0.208864	-1.959670	-1.328186	0.196861	0
4	-0.460639	1.057122	0.343618	-1.763040	0

Accuracy of Random Forest on test data: 53.00%

### Assignment 7: Write a program to implementing Boosting using AdaBoost

```
# Import necessary libraries
```

```
import pandas as pd
```

```
from sklearn.model_selection import train_test_split
```

```
from sklearn.ensemble import AdaBoostClassifier
```

```
from sklearn.tree import DecisionTreeClassifier
```

```
from sklearn.metrics import accuracy_score
```

```
from sklearn.preprocessing import LabelEncoder
```

```
# Load CSV file (replace 'your_dataset.csv' with your actual dataset file path)
```

```
df = pd.read_csv('log.csv')
```

```
# Check the first few rows of the dataset
```

```
print(df.head())
```

```
# Step 1: Handle missing values (if any)
```

```
# Example: Drop rows with missing values (you can also fill with the mean or median)
```

```
df = df.dropna()
```

```
# Step 2: Separate features (X) and target (y)
```

```
# Assuming the last column is the target variable
```

```

X = df.iloc[:, :-1].values # All rows, all columns except the last one (features)
y = df.iloc[:, -1].values # Last column is the target

# Step 3: Encode the target variable if it's categorical
label_encoder = LabelEncoder()
y = label_encoder.fit_transform(y)

# Step 4: Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)

# Step 5: Initialize and train the AdaBoost classifier with a DecisionTree as the base estimator
# DecisionTreeClassifier with max_depth=1 is used to create a weak learner (stump)
base_estimator = DecisionTreeClassifier(max_depth=1)

# Update: Use 'estimator' instead of 'base_estimator'
adaboost_classifier = AdaBoostClassifier(estimator=base_estimator, n_estimators=50,
random_state=42)
adaboost_classifier.fit(X_train, y_train)

# Step 6: Make predictions on the test set
y_pred = adaboost_classifier.predict(X_test)

# Step 7: Evaluate the model's performance
accuracy = accuracy_score(y_test, y_pred)
print(f'Accuracy of AdaBoost on test data: {accuracy * 100:.2f}%')

```

### Output:

	feature_0	feature_1	feature_2	feature_3	feature_4	feature_5	\
0	0.496714	-0.138264	0.647689	1.523030	-0.234153	-0.234137	
1	-0.463418	-0.465730	0.241962	-1.913280	-1.724918	-0.562288	
2	1.465649	-0.225776	0.067528	-1.424748	-0.544383	0.110923	
3	-0.601707	1.852278	-0.013497	-1.057711	0.822545	-1.220844	
4	0.738467	0.171368	-0.115648	-0.301104	-1.478522	-0.719844	

	feature_6	feature_7	feature_8	feature_9	target
0	1.579213	0.767435	-0.469474	0.542560	1
1	-1.012831	0.314247	-0.908024	-1.412304	1
2	-1.150994	0.375698	-0.600639	-0.291694	0
3	0.208864	-1.959670	-1.328186	0.196861	0
4	-0.460639	1.057122	0.343618	-1.763040	0

Accuracy of AdaBoost on test data: 49.00%

### Assignment 8: Write a program to implementing SVM for classification tasks.

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.svm import SVC
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix

# Load dataset from CSV file (replace 'your_dataset.csv' with your actual file path)
df = pd.read_csv('pca.csv')

# Check the first few rows of the dataset
print(df.head())

# Assuming the last column is the target variable (classification labels)
X = df.iloc[:, :-1] # Select all rows and all columns except the last one for features
y = df.iloc[:, -1] # Select the last column for the target (labels)

# Step 1: Split the dataset into training and testing sets (80% training, 20% testing)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Step 2: Standardize features using StandardScaler (important for SVM)
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train) # Fit on training data and transform it
X_test = scaler.transform(X_test)      # Use the same scaler to transform test data

# Step 3: Train the Support Vector Machine (SVM) model
svm_model = SVC(kernel='linear', random_state=42) # Using a linear kernel for simplicity
svm_model.fit(X_train, y_train) # Train the model on the training set

# Step 4: Make predictions using the trained SVM model
y_pred = svm_model.predict(X_test)

# Step 5: Evaluate the model's performance
accuracy = accuracy_score(y_test, y_pred)
conf_matrix = confusion_matrix(y_test, y_pred)
report = classification_report(y_test, y_pred)
```

```

# Print evaluation results
print(f'Accuracy: {accuracy:.4f}')
print('Confusion Matrix:')
print(conf_matrix)
print('Classification Report:')
print(report)

# Step 6: Optional - Visualize the decision boundaries (only works for 2D features)
# This is just a visualization example for datasets with two features
if X.shape[1] == 2: # Check if we have only two features for visualization
    # Create a mesh grid for plotting decision boundaries
    h = .02
    x_min, x_max = X_train[:, 0].min() - 1, X_train[:, 0].max() + 1
    y_min, y_max = X_train[:, 1].min() - 1, X_train[:, 1].max() + 1
    xx, yy = np.meshgrid(np.arange(x_min, x_max, h), np.arange(y_min, y_max, h))

    # Predict over the mesh grid
    Z = svm_model.predict(np.c_[xx.ravel(), yy.ravel()])
    Z = Z.reshape(xx.shape)

    # Plot decision boundary
    plt.contourf(xx, yy, Z, alpha=0.8)
    plt.scatter(X_train[:, 0], X_train[:, 1], c=y_train, edgecolors='k', marker='o',
cmap=plt.cm.Paired)
    plt.title('SVM Decision Boundary with Linear Kernel')
    plt.xlabel('Feature 1')
    plt.ylabel('Feature 2')
    plt.show()

```

### Output:

	Feature1	Feature2	Target
0	2.5	3.1	0
1	1.2	2.3	0
2	3.4	4.2	1
3	2.1	3.0	1
4	3.0	3.5	1

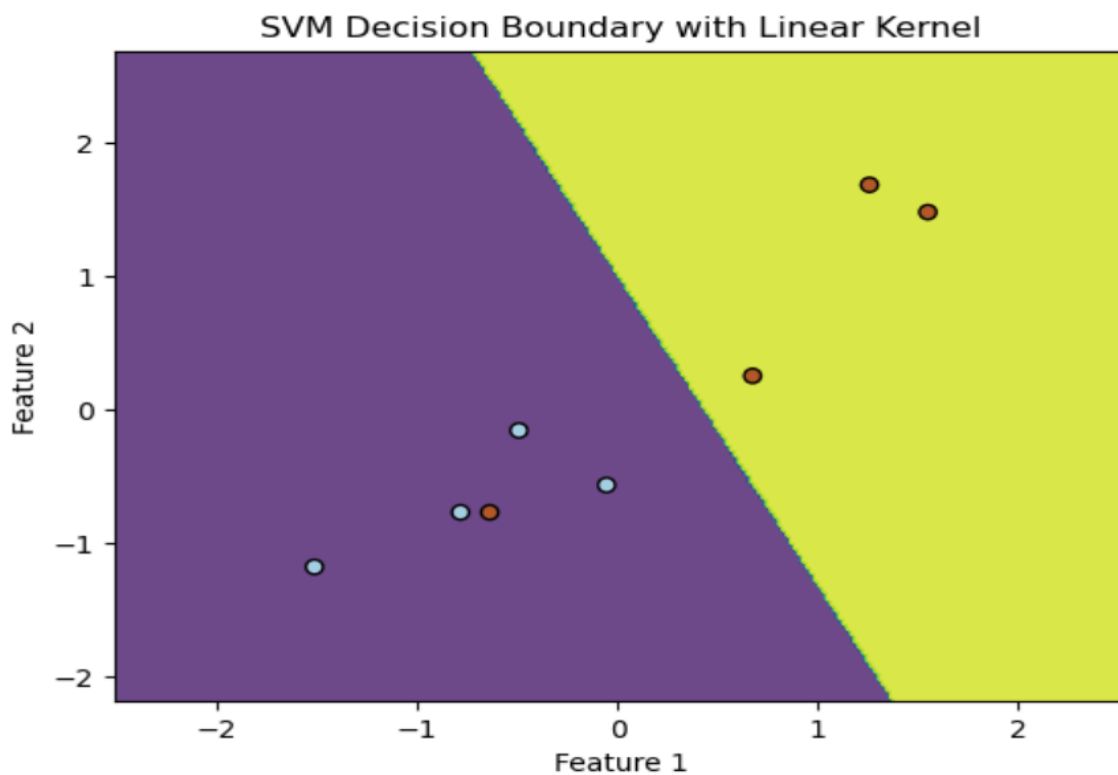
Accuracy: 1.0000

Confusion Matrix:

```
[[1 0]
 [0 1]]
```

Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	1
1	1.00	1.00	1.00	1
accuracy			1.00	2
macro avg	1.00	1.00	1.00	2
weighted avg	1.00	1.00	1.00	2



**Assignment 9: Write a program to implementing a simple neural network using TensorFlow/Keras.**

```
import pandas as pd
import numpy as np
import tensorflow as tf
import tensorflow as tf
from tensorflow.keras.models import Sequential
```

```

from tensorflow.keras.layers import Dense, Input
import matplotlib.pyplot as plt

df = pd.read_csv('diabetes.csv')
df.head()

print ('Number of Rows :', df.shape[0])
print ('Number of Columns :', df.shape[1])
print ('Number of Patients with outcome 1 :', df.Outcome.sum())
print ('Event Rate :', round(df.Outcome.mean()*100,2) ,'%')
df.describe()

from sklearn.model_selection import train_test_split
X = df.to_numpy()[:,0:8]
Y = df.to_numpy()[:,8]
seed = 42
X_train, X_test, y_train, y_test = train_test_split(X, Y, test_size = 0.25, random_state = seed)
print (f'Shape of Train Data : {X_train.shape}')
print (f'Shape of Test Data : {X_test.shape}')

model = Sequential([
    Input(shape=(8,)), # Define the input shape using the new `shape` argument
    Dense(24, activation='relu'),
    Dense(12, activation='relu'),
    Dense(1, activation='sigmoid'),
])

# Compile the model (optional, but necessary for training)
model.compile(optimizer = 'adam', loss = 'binary_crossentropy', metrics = ['accuracy'])

# Summary of the model
#model.summary()
model.summary()
history = model.fit(X_train, y_train, epochs=150, batch_size=32, verbose = 1)

# Plotting loss
plt.plot(history.history['loss'])
plt.title('Binary Cross Entropy Loss on Train dataset')
plt.ylabel('loss')
plt.xlabel('epoch')

```



```
plt.show()
```

```
# Plotting accuracy metric
plt.plot(history.history['accuracy'])
plt.title('Accuracy on the train dataset')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.show()
```

### Output:

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
0	6	148	72	35	0	33.6	0.627	50	1
1	1	85	66	29	0	26.6	0.351	31	0
2	8	183	64	0	0	23.3	0.672	32	1
3	1	89	66	23	94	28.1	0.167	21	0
4	0	137	40	35	168	43.1	2.288	33	1

Number of Rows : 768  
Number of Columns : 9  
Number of Patients with outcome 1 : 268  
Event Rate : 34.9 %  
Shape of Train Data : (576, 8)  
Shape of Test Data : (192, 8)

Model: "sequential\_6"

Layer (type)	Output Shape	Param #
dense_18 (Dense)	(None, 24)	216
dense_19 (Dense)	(None, 12)	300
dense_20 (Dense)	(None, 1)	13

Total params: 529 (2.07 KB)

Trainable params: 529 (2.07 KB)

Non-trainable params: 0 (0.00 B)

Epoch 1/150

18/18  1s 3ms/step - accuracy: 0.5895 - loss: 4.9180

Epoch 2/150

18/18 ————— 0s 3ms/step - accuracy: 0.5057 - loss:  
1.5014  
Epoch 3/150  
18/18 ————— 0s 3ms/step - accuracy: 0.6015 - loss:  
1.1479  
Epoch 4/150  
18/18 ————— 0s 3ms/step - accuracy: 0.6004 - loss:  
1.0280  
Epoch 5/150  
18/18 ————— 0s 2ms/step - accuracy: 0.6393 - loss:  
0.8963  
Epoch 6/150  
18/18 ————— 0s 2ms/step - accuracy: 0.6242 - loss:  
0.8435  
Epoch 7/150  
18/18 ————— 0s 2ms/step - accuracy: 0.6523 - loss:  
0.7196  
Epoch 8/150  
18/18 ————— 0s 2ms/step - accuracy: 0.6367 - loss:  
0.7654  
Epoch 9/150  
18/18 ————— 0s 2ms/step - accuracy: 0.6541 - loss:  
0.7686  
Epoch 10/150  
18/18 ————— 0s 2ms/step - accuracy: 0.6621 - loss:  
0.6817  
Epoch 11/150  
18/18 ————— 0s 5ms/step - accuracy: 0.6425 - loss:  
0.7043  
Epoch 12/150  
18/18 ————— 0s 2ms/step - accuracy: 0.7117 - loss:  
0.6386  
Epoch 13/150  
18/18 ————— 0s 3ms/step - accuracy: 0.6487 - loss:  
0.6587  
Epoch 14/150  
18/18 ————— 0s 2ms/step - accuracy: 0.6393 - loss:  
0.6852  
Epoch 15/150

18/18 ————— 0s 2ms/step - accuracy: 0.6932 - loss:  
0.6562  
Epoch 16/150  
18/18 ————— 0s 2ms/step - accuracy: 0.6278 - loss:  
0.6890  
Epoch 17/150  
18/18 ————— 0s 2ms/step - accuracy: 0.6813 - loss:  
0.6323  
Epoch 18/150  
18/18 ————— 0s 3ms/step - accuracy: 0.6853 - loss:  
0.6206  
Epoch 19/150  
18/18 ————— 0s 2ms/step - accuracy: 0.7143 - loss:  
0.5713  
Epoch 20/150  
18/18 ————— 0s 2ms/step - accuracy: 0.6995 - loss:  
0.5817  
Epoch 21/150  
18/18 ————— 0s 2ms/step - accuracy: 0.6969 - loss:  
0.6234  
Epoch 22/150  
18/18 ————— 0s 2ms/step - accuracy: 0.7152 - loss:  
0.5890  
Epoch 23/150  
18/18 ————— 0s 2ms/step - accuracy: 0.7173 - loss:  
0.5726  
Epoch 24/150  
18/18 ————— 0s 2ms/step - accuracy: 0.6975 - loss:  
0.5961  
Epoch 25/150  
18/18 ————— 0s 2ms/step - accuracy: 0.6997 - loss:  
0.6367  
Epoch 26/150  
18/18 ————— 0s 2ms/step - accuracy: 0.7037 - loss:  
0.6078  
Epoch 27/150  
18/18 ————— 0s 2ms/step - accuracy: 0.7107 - loss:  
0.5834  
Epoch 28/150

18/18 ————— 0s 2ms/step - accuracy: 0.7045 - loss:  
0.5672  
Epoch 29/150  
18/18 ————— 0s 2ms/step - accuracy: 0.7128 - loss:  
0.5608  
Epoch 30/150  
18/18 ————— 0s 3ms/step - accuracy: 0.7039 - loss:  
0.5908  
Epoch 31/150  
18/18 ————— 0s 3ms/step - accuracy: 0.7340 - loss:  
0.5645  
Epoch 32/150  
18/18 ————— 0s 2ms/step - accuracy: 0.7118 - loss:  
0.5859  
Epoch 33/150  
18/18 ————— 0s 2ms/step - accuracy: 0.7159 - loss:  
0.5662  
Epoch 34/150  
18/18 ————— 0s 2ms/step - accuracy: 0.6892 - loss:  
0.5822  
Epoch 35/150  
18/18 ————— 0s 2ms/step - accuracy: 0.7272 - loss:  
0.5440  
Epoch 36/150  
18/18 ————— 0s 3ms/step - accuracy: 0.7180 - loss:  
0.5606  
Epoch 37/150  
18/18 ————— 0s 2ms/step - accuracy: 0.7144 - loss:  
0.5587  
Epoch 38/150  
18/18 ————— 0s 2ms/step - accuracy: 0.7395 - loss:  
0.5620  
Epoch 39/150  
18/18 ————— 0s 3ms/step - accuracy: 0.7001 - loss:  
0.5776  
Epoch 40/150  
18/18 ————— 0s 2ms/step - accuracy: 0.6902 - loss:  
0.5962  
Epoch 41/150

18/18 ————— 0s 3ms/step - accuracy: 0.7072 - loss:  
0.5891  
Epoch 42/150  
18/18 ————— 0s 3ms/step - accuracy: 0.7389 - loss:  
0.5666  
Epoch 43/150  
18/18 ————— 0s 3ms/step - accuracy: 0.6933 - loss:  
0.6019  
Epoch 44/150  
18/18 ————— 0s 3ms/step - accuracy: 0.7406 - loss:  
0.5502  
Epoch 45/150  
18/18 ————— 0s 2ms/step - accuracy: 0.7194 - loss:  
0.5679  
Epoch 46/150  
18/18 ————— 0s 2ms/step - accuracy: 0.7061 - loss:  
0.5431  
Epoch 47/150  
18/18 ————— 0s 2ms/step - accuracy: 0.7438 - loss:  
0.5202  
Epoch 48/150  
18/18 ————— 0s 2ms/step - accuracy: 0.7346 - loss:  
0.5445  
Epoch 49/150  
18/18 ————— 0s 2ms/step - accuracy: 0.6963 - loss:  
0.5831  
Epoch 50/150  
18/18 ————— 0s 2ms/step - accuracy: 0.7246 - loss:  
0.5752  
Epoch 51/150  
18/18 ————— 0s 2ms/step - accuracy: 0.7290 - loss:  
0.5586  
Epoch 52/150  
18/18 ————— 0s 2ms/step - accuracy: 0.6933 - loss:  
0.5897  
Epoch 53/150  
18/18 ————— 0s 2ms/step - accuracy: 0.7588 - loss:  
0.5280  
Epoch 54/150

18/18 ————— 0s 2ms/step - accuracy: 0.7153 - loss:  
0.5653  
Epoch 55/150  
18/18 ————— 0s 2ms/step - accuracy: 0.7419 - loss:  
0.5464  
Epoch 56/150  
18/18 ————— 0s 2ms/step - accuracy: 0.7117 - loss:  
0.5628  
Epoch 57/150  
18/18 ————— 0s 2ms/step - accuracy: 0.7576 - loss:  
0.5011  
Epoch 58/150  
18/18 ————— 0s 2ms/step - accuracy: 0.7324 - loss:  
0.5320  
Epoch 59/150  
18/18 ————— 0s 2ms/step - accuracy: 0.7703 - loss:  
0.4985  
Epoch 60/150  
18/18 ————— 0s 3ms/step - accuracy: 0.7558 - loss:  
0.5444  
Epoch 61/150  
18/18 ————— 0s 2ms/step - accuracy: 0.7404 - loss:  
0.5359  
Epoch 62/150  
18/18 ————— 0s 2ms/step - accuracy: 0.7296 - loss:  
0.5656  
Epoch 63/150  
18/18 ————— 0s 2ms/step - accuracy: 0.7643 - loss:  
0.5179  
Epoch 64/150  
18/18 ————— 0s 2ms/step - accuracy: 0.7075 - loss:  
0.5770  
Epoch 65/150  
18/18 ————— 0s 2ms/step - accuracy: 0.7425 - loss:  
0.5180  
Epoch 66/150  
18/18 ————— 0s 2ms/step - accuracy: 0.7389 - loss:  
0.5451  
Epoch 67/150

18/18 ————— 0s 2ms/step - accuracy: 0.7635 - loss:  
0.5084  
Epoch 68/150  
18/18 ————— 0s 2ms/step - accuracy: 0.7268 - loss:  
0.5460  
Epoch 69/150  
18/18 ————— 0s 2ms/step - accuracy: 0.7467 - loss:  
0.5410  
Epoch 70/150  
18/18 ————— 0s 2ms/step - accuracy: 0.7459 - loss:  
0.5315  
Epoch 71/150  
18/18 ————— 0s 2ms/step - accuracy: 0.7889 - loss:  
0.5101  
Epoch 72/150  
18/18 ————— 0s 2ms/step - accuracy: 0.7333 - loss:  
0.5205  
Epoch 73/150  
18/18 ————— 0s 2ms/step - accuracy: 0.7744 - loss:  
0.5036  
Epoch 74/150  
18/18 ————— 0s 2ms/step - accuracy: 0.7172 - loss:  
0.5641  
Epoch 75/150  
18/18 ————— 0s 2ms/step - accuracy: 0.7345 - loss:  
0.5297  
Epoch 76/150  
18/18 ————— 0s 2ms/step - accuracy: 0.7540 - loss:  
0.5003  
Epoch 77/150  
18/18 ————— 0s 2ms/step - accuracy: 0.7523 - loss:  
0.5178  
Epoch 78/150  
18/18 ————— 0s 2ms/step - accuracy: 0.7509 - loss:  
0.5055  
Epoch 79/150  
18/18 ————— 0s 2ms/step - accuracy: 0.7389 - loss:  
0.5510  
Epoch 80/150

18/18 ————— 0s 2ms/step - accuracy: 0.7464 - loss:  
0.5330  
Epoch 81/150

18/18 ————— 0s 4ms/step - accuracy: 0.7178 - loss:  
0.5377  
Epoch 82/150

18/18 ————— 0s 2ms/step - accuracy: 0.7065 - loss:  
0.5838  
Epoch 83/150

18/18 ————— 0s 2ms/step - accuracy: 0.7326 - loss:  
0.5747  
Epoch 84/150

18/18 ————— 0s 2ms/step - accuracy: 0.7411 - loss:  
0.5718  
Epoch 85/150

18/18 ————— 0s 2ms/step - accuracy: 0.7313 - loss:  
0.6302  
Epoch 86/150

18/18 ————— 0s 2ms/step - accuracy: 0.7352 - loss:  
0.5402  
Epoch 87/150

18/18 ————— 0s 2ms/step - accuracy: 0.7502 - loss:  
0.5175  
Epoch 88/150

18/18 ————— 0s 2ms/step - accuracy: 0.7232 - loss:  
0.5524  
Epoch 89/150

18/18 ————— 0s 2ms/step - accuracy: 0.7249 - loss:  
0.5583  
Epoch 90/150

18/18 ————— 0s 2ms/step - accuracy: 0.7564 - loss:  
0.5251  
Epoch 91/150

18/18 ————— 0s 2ms/step - accuracy: 0.7525 - loss:  
0.4848  
Epoch 92/150

18/18 ————— 0s 2ms/step - accuracy: 0.7200 - loss:  
0.5333  
Epoch 93/150



18/18 ————— 0s 2ms/step - accuracy: 0.7700 - loss:  
0.4992  
Epoch 94/150  
18/18 ————— 0s 2ms/step - accuracy: 0.7563 - loss:  
0.5124  
Epoch 95/150  
18/18 ————— 0s 5ms/step - accuracy: 0.7918 - loss:  
0.4565  
Epoch 96/150  
18/18 ————— 0s 2ms/step - accuracy: 0.7253 - loss:  
0.5327  
Epoch 97/150  
18/18 ————— 0s 2ms/step - accuracy: 0.7681 - loss:  
0.4731  
Epoch 98/150  
18/18 ————— 0s 2ms/step - accuracy: 0.7563 - loss:  
0.5068  
Epoch 99/150  
18/18 ————— 0s 2ms/step - accuracy: 0.7770 - loss:  
0.4929  
Epoch 100/150  
18/18 ————— 0s 2ms/step - accuracy: 0.7342 - loss:  
0.5462  
Epoch 101/150  
18/18 ————— 0s 2ms/step - accuracy: 0.7210 - loss:  
0.5485  
Epoch 102/150  
18/18 ————— 0s 2ms/step - accuracy: 0.7474 - loss:  
0.5126  
Epoch 103/150  
18/18 ————— 0s 2ms/step - accuracy: 0.7680 - loss:  
0.5199  
Epoch 104/150  
18/18 ————— 0s 2ms/step - accuracy: 0.7530 - loss:  
0.5064  
Epoch 105/150  
18/18 ————— 0s 2ms/step - accuracy: 0.7746 - loss:  
0.5026  
Epoch 106/150

18/18 ————— 0s 2ms/step - accuracy: 0.7992 - loss:  
0.4694  
Epoch 107/150  
18/18 ————— 0s 2ms/step - accuracy: 0.7494 - loss:  
0.5140  
Epoch 108/150  
18/18 ————— 0s 2ms/step - accuracy: 0.7442 - loss:  
0.5407  
Epoch 109/150  
18/18 ————— 0s 2ms/step - accuracy: 0.7394 - loss:  
0.5266  
Epoch 110/150  
18/18 ————— 0s 4ms/step - accuracy: 0.7719 - loss:  
0.4795  
Epoch 111/150  
18/18 ————— 0s 2ms/step - accuracy: 0.7351 - loss:  
0.5200  
Epoch 112/150  
18/18 ————— 0s 2ms/step - accuracy: 0.7625 - loss:  
0.4970  
Epoch 113/150  
18/18 ————— 0s 2ms/step - accuracy: 0.7815 - loss:  
0.4912  
Epoch 114/150  
18/18 ————— 0s 2ms/step - accuracy: 0.7550 - loss:  
0.4981  
Epoch 115/150  
18/18 ————— 0s 2ms/step - accuracy: 0.7752 - loss:  
0.4759  
Epoch 116/150  
18/18 ————— 0s 2ms/step - accuracy: 0.7665 - loss:  
0.4901  
Epoch 117/150  
18/18 ————— 0s 2ms/step - accuracy: 0.7712 - loss:  
0.4704  
Epoch 118/150  
18/18 ————— 0s 2ms/step - accuracy: 0.7220 - loss:  
0.5439  
Epoch 119/150

18/18 ————— 0s 2ms/step - accuracy: 0.7663 - loss:  
0.4747  
Epoch 120/150  
18/18 ————— 0s 2ms/step - accuracy: 0.7419 - loss:  
0.5171  
Epoch 121/150  
18/18 ————— 0s 2ms/step - accuracy: 0.7474 - loss:  
0.5315  
Epoch 122/150  
18/18 ————— 0s 2ms/step - accuracy: 0.7596 - loss:  
0.5100  
Epoch 123/150  
18/18 ————— 0s 2ms/step - accuracy: 0.7611 - loss:  
0.4954  
Epoch 124/150  
18/18 ————— 0s 4ms/step - accuracy: 0.7780 - loss:  
0.4855  
Epoch 125/150  
18/18 ————— 0s 2ms/step - accuracy: 0.7708 - loss:  
0.4956  
Epoch 126/150  
18/18 ————— 0s 2ms/step - accuracy: 0.7892 - loss:  
0.4946  
Epoch 127/150  
18/18 ————— 0s 2ms/step - accuracy: 0.7564 - loss:  
0.4991  
Epoch 128/150  
18/18 ————— 0s 2ms/step - accuracy: 0.7740 - loss:  
0.5154  
Epoch 129/150  
18/18 ————— 0s 2ms/step - accuracy: 0.7924 - loss:  
0.4877  
Epoch 130/150  
18/18 ————— 0s 2ms/step - accuracy: 0.7559 - loss:  
0.4971  
Epoch 131/150  
18/18 ————— 0s 2ms/step - accuracy: 0.7745 - loss:  
0.4818  
Epoch 132/150

18/18 ————— 0s 2ms/step - accuracy: 0.7583 - loss:  
0.4932  
Epoch 133/150  
18/18 ————— 0s 2ms/step - accuracy: 0.7677 - loss:  
0.4927  
Epoch 134/150  
18/18 ————— 0s 2ms/step - accuracy: 0.7872 - loss:  
0.4741  
Epoch 135/150  
18/18 ————— 0s 2ms/step - accuracy: 0.7663 - loss:  
0.4600  
Epoch 136/150  
18/18 ————— 0s 4ms/step - accuracy: 0.7785 - loss:  
0.4928  
Epoch 137/150  
18/18 ————— 0s 3ms/step - accuracy: 0.7636 - loss:  
0.5001  
Epoch 138/150  
18/18 ————— 0s 2ms/step - accuracy: 0.7424 - loss:  
0.5203  
Epoch 139/150  
18/18 ————— 0s 2ms/step - accuracy: 0.7641 - loss:  
0.4943  
Epoch 140/150  
18/18 ————— 0s 2ms/step - accuracy: 0.7559 - loss:  
0.4887  
Epoch 141/150  
18/18 ————— 0s 2ms/step - accuracy: 0.7780 - loss:  
0.4835  
Epoch 142/150  
18/18 ————— 0s 2ms/step - accuracy: 0.7715 - loss:  
0.5056  
Epoch 143/150  
18/18 ————— 0s 2ms/step - accuracy: 0.7542 - loss:  
0.5170  
Epoch 144/150  
18/18 ————— 0s 2ms/step - accuracy: 0.7890 - loss:  
0.4876  
Epoch 145/150

18/18 ————— 0s 2ms/step - accuracy: 0.8027 - loss:  
0.4567

Epoch 146/150

18/18 ————— 0s 2ms/step - accuracy: 0.7647 - loss:  
0.5335

Epoch 147/150

18/18 ————— 0s 2ms/step - accuracy: 0.7290 - loss:  
0.5359

Epoch 148/150


18/18 ————— 0s 3ms/step - accuracy: 0.7641 - loss:  
0.5116


Epoch 149/150

18/18 ————— 0s 2ms/step - accuracy: 0.7540 - loss:  
0.5059

Epoch 150/150

18/18 ————— 0s 2ms/step - accuracy: 0.7521 - loss:  
0.5285





**Assignment 10: Write a program to implementing with big data concepts using sample datasets & Setting up a Hadoop environment.**

#### # Install Java

```
!sudo apt update
!sudo apt install openjdk-8-jdk
```

#### # Download and extract Hadoop

```
!wget http://apache.mirrors.lucidnetworks.net/hadoop/common/hadoop-3.3.1/hadoop-3.3.1.tar.gz
!tar -xzf hadoop-3.3.1.tar.gz
!mv hadoop-3.3.1 /usr/local/hadoop
```

#### # Sample dataset (you can imagine it as a text file with large data)

```
dataset = """
Hadoop is a framework for processing large datasets.
It is used for distributed storage and distributed computing.
Hadoop is part of the Big Data ecosystem.
Hadoop helps process Big Data.
"""
```

```

# Save dataset to a file (simulating a big text file)
with open('/content/dataset.txt', 'w') as f:
    f.write(dataset)

from pyspark.sql import SparkSession

# Initialize Spark session
spark = SparkSession.builder.appName('WordCount').getOrCreate()

# Load the dataset into an RDD (Resilient Distributed Dataset)
rdd = spark.sparkContext.textFile('/content/dataset.txt')

# Perform word count
word_counts = rdd.flatMap(lambda line: line.split()) \
    .map(lambda word: (word.lower(), 1)) \
    .reduceByKey(lambda x, y: x + y)

# Collect and print the results
for word, count in word_counts.collect():
    print(f'{word}: {count}')

# Stop the Spark session
spark.stop()

```

### Output:

```

hadoop: 3
framework: 1
for: 2
large: 1
it: 1
used: 1
distributed: 2
storage: 1
and: 1
part: 1
of: 1
big: 2
ecosystem.: 1
helps: 1
data.: 1
is: 3
a: 1
processing: 1
datasets.: 1
computing.: 1
the: 1
data: 1
process: 1

```

## Assignment 11: Write a program to implementing CRUD operations in MongoDB

```
pip install pymongo
```

```
from pymongo import MongoClient
```

```
# Connect to MongoDB server (default localhost:27017)
```

```
client = MongoClient("mongodb://localhost:27017/")
```

```
# Use the 'mydatabase' database and 'users' collection
```

```
db = client['mydatabase']
```

```
collection = db['users']
```

```
# CREATE operation: Insert a document
```

```
user_data = {
```

```
    'name': 'John Doe',
```

```
    'age': 30,
```

```
    'email': 'john.doe@example.com'
```

```
}
```

```
result = collection.insert_one(user_data)
```

```
print(f"Document inserted with ID: {result.inserted_id}")
```

```
# READ operation: Find a single document by name
```

```
user = collection.find_one({"name": "John Doe"})
```

```
print("Found user:", user)
```

```
# UPDATE operation: Update the user's age
```

```
update_result = collection.update_one(
```

```
    {"name": "John Doe"},
```

```
    {"$set": {"age": 31}}
```

```
)
```

```
print(f"Documents matched: {update_result.matched_count}, Documents modified:
```

```
{update_result.modified_count}")
```

```
# DELETE operation: Delete a user by name
```

```
delete_result = collection.delete_one({"name": "John Doe"})
```

```
print(f"Documents deleted: {delete_result.deleted_count}")
```



## Output:

```
Collecting pymongo
  Downloading pymongo-4.11.2-cp312-cp312-win_amd64.whl.metadata (22 kB)
Collecting dnspython<3.0.0,>=1.16.0 (from pymongo)
  Downloading dnspython-2.7.0-py3-none-any.whl.metadata (5.8 kB)
Downloaded pymongo-4.11.2-cp312-cp312-win_amd64.whl (882 kB)
----- 0.0/882.2 kB ? eta -:--:--
----- 882.2/882.2 kB 19.4 MB/s eta 0:00:00
Downloaded dnspython-2.7.0-py3-none-any.whl (313 kB)
Installing collected packages: dnspython, pymongo
Successfully installed dnspython-2.7.0 pymongo-4.11.2
Note: you may need to restart the kernel to use updated packages.
```

Document inserted with ID: 67d270cb7c6068e82f03a444

Found user: {'\_id': ObjectId('67d270cb7c6068e82f03a444'), 'name': 'John Doe', 'age': 30, 'email': 'john.doe@example.com'}

Documents matched: 1, Documents modified: 1

Documents deleted: 1

## Assignment 12: Write a program to implementing with NLTK: Tokenization, stemming, and lemmatization

```
pip install nltk
```

```
import nltk
```

```
# Download the 'punkt' tokenizer
```

```
nltk.download('punkt')
```

```
import nltk
```

```
# Download the 'punkt_tab' resource, which is required for tokenization
```

```
nltk.download('punkt_tab')
```

```
# Download other necessary resources for lemmatization and stop words
```

```
nltk.download('wordnet')
```

```
nltk.download('stopwords')
```

```
from nltk.tokenize import word_tokenize
```

```
from nltk.stem import PorterStemmer
```

```

from nltk.stem import WordNetLemmatizer

# Sample text for demonstration
text = "NLTK is a great toolkit for Natural Language Processing. Tokenization, Stemming, and
Lemmatization are important tasks."

# Tokenization: Split text into words
tokens = word_tokenize(text)
print("Tokens:", tokens)

# Stemming: Reduce words to their root form using Porter Stemmer
stemmer = PorterStemmer()
stemmed_words = [stemmer.stem(word) for word in tokens]
print("Stemmed words:", stemmed_words)

# Lemmatization: Reduce words to their base form using WordNet Lemmatizer
lemmatizer = WordNetLemmatizer()
lemmatized_words = [lemmatizer.lemmatize(word) for word in tokens]
print("Lemmatized words (default pos=noun):", lemmatized_words)
# Optional: Lemmatization with POS tagging (verbs, adjectives, etc.)
lemmatized_verbs = [lemmatizer.lemmatize(word, pos='v') for word in tokens]
print("Lemmatized words (as verbs):", lemmatized_verbs)

```

### Output:

```

Requirement already satisfied: nltk in c:\users\imrd\anaconda3\lib\site-packages (3.9.1)
Requirement already satisfied: click in c:\users\imrd\anaconda3\lib\site-packages (from nltk) (8.1.7)
Requirement already satisfied: joblib in c:\users\imrd\anaconda3\lib\site-packages (from nltk) (1.4.2)
Requirement already satisfied: regex>=2021.8.3 in c:\users\imrd\anaconda3\lib\site-packages (from nltk) (2024.9.11)
Requirement already satisfied: tqdm in c:\users\imrd\anaconda3\lib\site-packages (from nltk) (4.66.5)
Requirement already satisfied: colorama in c:\users\imrd\anaconda3\lib\site-packages (from click->nltk) (0.4.6)
Note: you may need to restart the kernel to use updated packages.

```

```

[nltk_data] Downloading package punkt to
[nltk_data] C:\Users\IMRD\AppData\Roaming\nltk_data...
[nltk_data] Package punkt is already up-to-date!

```

True

[nltk\_data] Downloading package punkt\_tab to  
[nltk\_data] C:\Users\IMRD\AppData\Roaming\nltk\_data...  
[nltk\_data] Unzipping tokenizers\punkt\_tab.zip.  
[nltk\_data] Downloading package wordnet to  
[nltk\_data] C:\Users\IMRD\AppData\Roaming\nltk\_data...  
[nltk\_data] Package wordnet is already up-to-date!  
[nltk\_data] Downloading package stopwords to  
[nltk\_data] C:\Users\IMRD\AppData\Roaming\nltk\_data...  
[nltk\_data] Package stopwords is already up-to-date!

Tokens: ['NLTK', 'is', 'a', 'great', 'toolkit', 'for', 'Natural', 'Language', 'Processing', '.',  
'Tokenization', '.', 'Stemming', '.', 'and', 'Lemmatization', 'are', 'important', 'tasks', '.']  
Stemmed words: ['nltk', 'is', 'a', 'great', 'toolkit', 'for', 'natur', 'languag', 'process', '.', 'token', '.',  
'stem', '.', 'and', 'lemmat', 'are', 'import', 'task', '.']  
Lemmatized words (default pos=noun): ['NLTK', 'is', 'a', 'great', 'toolkit', 'for', 'Natural',  
'Language', 'Processing', '.', 'Tokenization', '.', 'Stemming', '.', 'and', 'Lemmatization', 'are',  
'important', 'task', '.']  
Lemmatized words (as verbs): ['NLTK', 'be', 'a', 'great', 'toolkit', 'for', 'Natural', 'Language',  
'Processing', '.', 'Tokenization', '.', 'Stemming', '.', 'and', 'Lemmatization', 'be', 'important', 'task', '.']